

EGRD: ASSESSING RESEARCH IDEA NOVELTY THROUGH EVIDENCE-GROUNDED RESIDUAL DECOMPOSITION

Anonymous authors

Paper under double-blind review

ABSTRACT

Language models now propose research ideas faster than anyone can vet them, which makes assessing novelty a practical bottleneck. Judging novelty requires identifying what prior work already covers and evaluating the contribution that remains. Overall similarity provides an incomplete account: an idea’s components may be covered across several papers, while their combination may still introduce a substantive change. We propose Evidence-Grounded Residual Decomposition (EGRD), which assesses novelty relative to an explicit set of temporally admissible prior works, making the basis of each judgment inspectable and allowing the effect of changing the evidence to be examined. EGRD estimates how much of the target those priors jointly cover, decomposes the target into latent semantic units, and learns a support matrix aligning these units with the evidence. It models two complementary residuals: unit residuals capture content with limited historical support, while interaction residuals capture unit combinations that no single prior supports. These residuals, together with coverage and contextual representations, predict an ordered novelty grade. A conservative blockwise aggregation rule lets a model trained with a small evidence budget read larger evidence sets at inference time without retraining. We further construct IDEADELTA, a dataset spanning graph machine learning, computer vision, and time series, pairing each idea with temporally admissible priors annotated for coverage, residual contribution, and novelty. EGRD recovers the annotation scheme consistently across the three domains, and ablations show that its predictions depend on the supplied evidence and on the residual representations rather than on the target idea alone. Transfer to expert judgments remains sensitive to where the grade boundaries are placed.

1 INTRODUCTION

Language models now propose research ideas faster than anyone can vet them, and deciding which of those ideas is new is the step that does not scale. Novelty is not a monotone function of textual similarity to any single prior work. We therefore study *evidence-conditioned* novelty assessment: grading a target idea against an explicit set of prior works published before it, which makes the basis of each judgment inspectable and turns the evidence into a variable we can intervene on.

Most prior work takes a paper as the unit of assessment, from citation-pattern measures (Uzzi et al., 2013) to document-level embeddings (Cohan et al., 2020; Ostendorff et al., 2022) to systems that review a submission as a whole (Zhu et al., 2025). But a paper bundles contributions of uneven novelty, and one score over the bundle does not say which contribution it describes. Systems that do operate on ideas assemble the judgment as a language-model workflow: retrieve related work, then prompt a model to compare against it (Wang et al., 2024; da Silva et al., 2025; Afzal et al., 2026). Here the evidence is selected by the system being evaluated, so no verdict can be attributed to the retrieval or to the comparison, and the comparison itself is elicited by a prompt rather than learned against supervision. Evaluating this arrangement against expert judgments, Schopf & Färber (2026) find that the reasoning such models produce mirrors human rationales, yet that alignment does not translate into accurate grades.

Such a workflow also compresses the comparison into a single score, and three distinctions do not survive. Coverage can be *dispersed*: several papers may jointly account for an idea that no single paper resembles. Ideas at equal coverage differ in *what remains*: one leaves an implementation detail unexplained, another introduces a new mechanism. And individually covered components may form a *combination* that no single prior work puts together. Each distinction asks about a particular part of the idea and the particular prior work responsible for it, and a score that discards that correspondence cannot be post-processed to recover it. Assessing novelty calls instead for a representation that keeps it.

We propose Evidence-Grounded Residual Decomposition (EGRD), which decomposes a target against ranked historical evidence into what that evidence accounts for and what it leaves, then predicts an ordered novelty grade from the two. Because existing benchmarks do not pair ideas with the evidence admissible when they were proposed, we also construct IDEADELTA to supply that pairing. Our contributions are:

- **IDEADELTA, a multi-domain evidence-paired dataset.** 21,111 ideas paired with temporally admissible priors and annotated at the idea and pair levels (92,876 pairs) across graph machine learning, computer vision, and time series. Supplied rather than retrieved, the evidence can be intervened on: permuting priors while holding targets fixed drops accuracy from 0.78 to 0.35, the majority-class band.
- **Unit-level coverage over dispersed evidence.** We model coverage at the level of latent semantic units, aligning each unit with each prior through a learned support matrix. Noisy-OR aggregation lets several partially relevant papers jointly cover a unit none covers alone; removing it costs 0.27 quadratic weighted kappa (QWK).
- **Typed residuals.** We split what the priors leave over into unit residuals (content no prior supports) and interaction residuals (combinations no single prior puts together). Both read off the same support matrix, so an unexplored contribution and an effective recombination are distinguished by the model rather than by assertion.
- **Evidence that the decomposition, not the backbone, carries the result.** EGRD reaches 0.909/0.882/0.898 QWK across the three domains against 0.579 for the best prompted system, whose scores vary more with the backbone than with the prompt chain. A block-wise rule extends the frozen checkpoint beyond its training evidence budget, leading every RINoBench (Schopf & Färber, 2026) baseline on QWK (0.2245 against 0.1469) and Macro-F1 (0.3407 against 0.2457); that margin is modest.

Appendix A situates EGRD against automated idea generation, prior novelty judges, and decomposed verification against evidence.

2 DATASET

Novelty can only be assessed relative to earlier work. Each evaluation example must therefore specify a target idea, the prior ideas available when it was proposed, and the contribution that remains after accounting for those priors. We construct IDEADELTA to provide these three elements.

2.1 TASK AND LABELS

The unit of annotation is a research *idea*: one innovation claim extracted from a paper. For example, a paper that introduces both a new objective and a new architecture contributes two ideas, which are evaluated separately. Let x denote a target idea published at time t_x . Its admissible priors are $\mathcal{P}_x = \{p_k \mid t_{p_k} < t_x\}$. This temporal constraint is essential: an idea published after the target cannot serve as evidence about the target’s novelty, regardless of its semantic similarity. We annotate each target–prior pair with one of four coverage levels and each target with an ordered novelty grade from LOW to HIGH.

Rather than assign novelty as a single overall impression, we derive it from four annotated quantities:

$$\text{raw}(x) = (1 - C_x) \cdot \Delta_x \cdot E_x \cdot (1 - \lambda S_x), \quad \lambda = 0.45, \quad (1)$$

where C_x measures how much of the target is jointly covered by the retained priors, Δ_x measures the substantive contribution that remains, E_x measures whether the retrieved evidence is sufficient

Table 1: Dataset composition. Counts of ideas and of target–prior pairs, and the idea-level novelty distribution. Pairs are formed within a source domain only.

	Graph ML	Computer vision	Time series	All
Ideas	5,225	8,291	7,595	21,111
Pairs	21,183	36,238	35,455	92,876
LOW	819	901	1,724	3,444
WEAK	1,240	1,711	1,614	4,565
MODERATE	2,247	2,777	2,173	7,197
HIGH	919	2,902	2,084	5,905

for assessment, and S_x measures how much of the claimed difference is merely surface-level. We convert the resulting score into the four ordered novelty grades. The rule assigns lower novelty when prior work already covers the target, when the remaining difference is not substantive, when the available evidence is insufficient, or when the difference consists mainly of renaming or other surface changes.

2.2 CONSTRUCTION

Paper selection and idea extraction. We collect arXiv papers from graph machine learning, computer vision, and time series. We prioritize papers that were subsequently published, especially at major venues. An LLM extracts innovation claims from the title, abstract, and introduction. A second pass removes statements that do not describe a substantive research contribution, including open-source announcements, bare state-of-the-art claims, and engineering deliverables.

Prior retrieval. For each target, nearest-neighbor search over idea embeddings retrieves candidate priors published before the target. An LLM adjudicates candidates in the uncertain similarity range. We also perform a cross-type retrieval pass so that a target can be linked to relevant prior work of another contribution type. All target–prior pairs remain within the same research domain.

Annotation. The same judge, `deepseek-v4-pro`, annotates all stages. It first assigns pairwise coverage labels, with ambiguous cases judged from their technical content. It then estimates joint coverage, the remaining substantive contribution, evidence sufficiency, and surface-level difference for each target. Equation (1) converts these quantities into the idea-level novelty grade. The labels are model-generated, and we manually inspect samples from each stage. Appendix B.1 describes the full pipeline.

2.3 COMPOSITION

IDEADELTA contains 21,111 annotated ideas and 92,876 annotated target–prior pairs, with target dates ranging from 2019-01 to 2026-09 (Table 1). Each idea records its source domain, and the *All* column is the union of the three domain subsets rather than an additional collection. We order the train, validation, and test splits by date, with cut points at 2025-02 and 2025-11. The resulting splits contain 59,269/16,436/17,171 pairs and 13,815/3,944/3,352 ideas.

Because the labels are model-generated, we further test whether they behave as novelty labels should. Appendix B.2 examines whether metadata can predict the labels, whether the retrieved prior content matters, and whether the assigned novelty changes appropriately when the evidence is modified.

3 METHODOLOGY

Given a target research idea and a set of prior ideas proposed before it, EGRD assesses the target’s novelty relative to the supplied evidence. The model proceeds in three stages, shown in Figure 1. Stage A assembles the historical evidence: it selects temporally admissible priors and summarizes how much of the target they already account for. Stage B decomposes the target against that evidence and represents the uncovered contribution through two complementary components, where




Figure 1: Overview of EGRD. (A) Only priors predating the target are admissible; the retained top- K evidence carries pair coverage features f_k , and a multiple-instance model predicts joint coverage C_J and its four-class distribution π^J . Stage A outputs enter the later stages as fixed inputs, marked $\text{sg}[\cdot]$. (B) M learned queries partition a target token bank into semantic units u_m ; a learned support matrix A_{km} aligns each unit with each prior, yielding the unit residual z_U , the covered pooling z_C , and—through co-support b_{ij} —the interaction residual z_I . (C) A gate modulates global context within residual fusion; the fused, covered, target and context representations form the shared representation feeding three novelty branches and the moderate-to-high refinement. The continuous score head supplies training supervision only.

unit residuals capture semantic content with insufficient historical support and interaction residuals capture potentially novel combinations of semantic units. Stage C predicts an ordered novelty grade from LOW to HIGH from these residual representations.

Throughout, $[\cdot; \cdot]$ denotes concatenation, $\odot$ elementwise multiplication, d the encoder dimension, σ the sigmoid function, and ε a small positive constant added for numerical stability. Padded priors are masked out, and $\mathcal{V}$ denotes the non-padding positions of the target.

3.1 STAGE A: HISTORICAL EVIDENCE INPUTS

Stage A assembles historical evidence and estimates how much of the target idea that evidence collectively covers. Given a target idea x with timestamp t_x , we restrict candidate priors to $\mathcal{P}_x = \{p_k \mid t_{p_k} < t_x\}$, where t_{p_k} denotes the timestamp of prior idea p_k . We retain the top- K ranked priors, or all candidates if fewer are available, to form the evidence set $\mathcal{E}_x \subseteq \mathcal{P}_x$ ($K = 5$ in our experiments). Each retained prior p_k carries a vector of pair coverage features $f_k \in \mathbb{R}^{11}$, listed in Appendix C.1.

To estimate coverage across the evidence set, we use a multiple-instance model (Ilse et al., 2018). A cross-encoder encodes each target–prior pair, and the model aggregates the pair representations through attention pooling and maximum pooling. The pooled representations are concatenated with evidence-score statistics to predict joint coverage C_J and a four-class coverage distribution π^J . Appendix C.1 specifies the pair-scoring rule and the joint model. The retained evidence, pair features, and coverage predictions serve as fixed inputs to Stages B and C, grounding the residual decomposition and providing coverage information for novelty prediction.

We summarize the evidence set in a global context vector z_G , a learned projection of twenty-one scalars—the scaled target timestamp, ten pair-score statistics over $\mathcal{E}_x$, and ten joint-coverage features (Appendix C.1)—concatenated with embeddings of two target facets (aspect and contribution type).

Apart from the timestamp and the facets, $\mathbf{z}_G$ is a re-encoding of the Stage A outputs, and Stage C consumes it as a summary of what the evidence collectively establishes.

3.2 STAGE B: EVIDENCE-GROUNDED RESIDUAL DECOMPOSITION

Stage B decomposes the target into semantic units and estimates their support from the retained evidence. It represents the unexplained contribution through two complementary components: unit residuals for content with limited historical support, and interaction residuals for combinations with limited joint support from any single prior.

Semantic decomposition. We use a residual cross-encoder initialized from the Stage A model and trained separately from it. For each target–prior pair, the encoder produces a pooled representation $\mathbf{v}_k$, which we augment with projected pair coverage features: $\tilde{\mathbf{v}}_k = \mathbf{v}_k + f_P(\mathbf{f}_k)$. To obtain a shared representation of the target, we average its contextual token states at matching positions across valid pairs, forming a token bank $\{\mathbf{h}_j\}_{j=1}^N$, where N is the number of target tokens. We then use M learned queries $\mathbf{q}_m$ to softly assign these tokens to semantic units (Locatello et al., 2020) ($M = 8$ in our experiments). Queries compete for each token through a softmax over units; the resulting assignments are then normalized over tokens within each unit:

$$B_{mj} = \text{softmax}_m \left(\frac{\mathbf{q}_m^\top \mathbf{h}_j}{\sqrt{d}} \right), \quad \alpha_{mj} = \frac{B_{mj}}{\sum_{j' \in \mathcal{V}} B_{mj'} + \varepsilon}, \quad \mathbf{u}_m = \sum_{j \in \mathcal{V}} \alpha_{mj} \mathbf{h}_j. \quad (2)$$

The units have no predefined semantic roles. A learned activity score $v_m = \sigma(\mathbf{w}_v^\top \mathbf{u}_m + b_v)$ controls each unit’s contribution, while regularization discourages redundant token assignments.

Evidence-aligned unit residuals. We estimate how strongly each prior p_k supports each unit $\mathbf{u}_m$ through a learned support matrix:

$$A_{km} = \sigma(f_A([\mathbf{u}_m; \tilde{\mathbf{v}}_k; \mathbf{u}_m \odot \tilde{\mathbf{v}}_k; |\mathbf{u}_m - \tilde{\mathbf{v}}_k|; \mathbf{f}_k])), \quad (3)$$

where f_A is an MLP. Support is estimated separately for every target–prior pair, so different priors can account for different units. For $n = |\mathcal{E}_x|$ retained priors, noisy-OR aggregates support across priors into unit coverage c_m . Its complement, weighted by unit activity, defines the residual weight r_m :

$$c_m = 1 - \prod_{k=1}^n (1 - A_{km}), \quad r_m = v_m(1 - c_m), \quad \mathbf{z}_U = \frac{1}{M} \sum_m r_m \mathbf{u}_m. \quad (4)$$

Noisy-OR lets several partially relevant priors jointly cover a unit that none of them covers alone. We use it as a differentiable aggregation rule, not as a claim that the priors are independent sources of support (Appendix C.2). An active unit with little historical support receives a large residual weight in $\mathbf{z}_U$. We normalize by the fixed number of units M rather than by $\sum_m r_m$, so that the pooled representation retains the total residual mass instead of dividing it away. We also pool the supported content as $\mathbf{z}_C = M^{-1} \sum_m v_m c_m \mathbf{u}_m$. To connect unit coverage to supervision, a learned, activity-weighted pooling of c_m reconstructs joint coverage $\hat{C}_U$, supervised by annotated coverage and the fixed Stage A prediction. The entries of the support matrix are latent strengths learned without unit-level alignment labels.

Interaction residuals. Units that are individually covered may still form a new combination. For each unit pair $i < j$, an MLP maps their concatenation, elementwise product, and absolute difference to an interaction representation $\mathbf{t}_{ij}$, and a sigmoid head predicts its strength g_{ij} . We measure historical co-support by the largest product of the two units’ support strengths within a single prior. The co-support score b_{ij} , interaction residual weight r_{ij} , and pooled representation are

$$b_{ij} = \max_k A_{ki} A_{kj}, \quad r_{ij} = g_{ij}(1 - b_{ij}), \quad \mathbf{z}_I = \sum_{i < j} \eta_{ij} r_{ij} \mathbf{t}_{ij}, \quad (5)$$

where the learned attention η_{ij} includes a $\log(r_{ij} + \varepsilon)$ bias that favors pairs with larger residual weights. When different priors support the two units but no single prior strongly supports both, their unit residuals can be small while their interaction residual remains large, provided that g_{ij} is high. An auxiliary classifier on $\mathbf{z}_I$ is trained with idea-level labels for nontrivial combinations. Co-support serves as a proxy for an existing combination: a prior that supports both units need not cover the specific way they are combined in the target. Appendix C.2 provides the detailed pooling equations.

3.3 STAGE C: NOVELTY PREDICTION

Stage C predicts an ordered novelty grade from the residual representations and the evidence context. To account for evidence availability and uncertainty, we condition residual fusion on the global context z_G defined in Section 3.1. A learned gate γ takes z_U, z_I, z_G , and the entropy of the Stage A coverage prediction as inputs and scales the context within the residual branch:

$$z_R = f_R([z_U; z_I; \gamma z_G]), \quad \mathbf{h} = \text{LayerNorm}([z_R; z_C; \tilde{h}_x; z_G]), \quad (6)$$

where $\tilde{h}_x$ is obtained by masked mean pooling over the target token bank. The shared representation $\mathbf{h}$ retains the covered content, the target representation, and a direct path for global context. The gate therefore controls context only within residual fusion.

Auxiliary heads predict substantive difference, surface change, evidence sufficiency, and residual type. These predictions, together with $\mathbf{h}$ and joint coverage, form the inputs to three novelty branches: a direct linear classifier, a cumulative ordinal head that models grade order, and a two-layer MLP classifier. Each branch produces a four-class distribution, and their convex mixture determines an initial grade. A moderate-to-high refinement upgrades an initial MODERATE prediction when a weighted combination of the direct and ordinal HIGH-class probabilities and a separate binary-head score exceeds a threshold. The mixture weights and refinement threshold are selected on development data and fixed for testing (Appendix C.3). A separate continuous novelty head provides auxiliary supervision during training; its output is not used in the inference-time grading rule.

3.4 TRAINING OBJECTIVE

We initialize the residual encoder from the Stage A joint-coverage checkpoint and jointly optimize it with the decomposition and prediction modules in Stages B and C, while keeping the Stage A models fixed. The objective combines classification of the novelty grade and the residual type, supervision of a family of annotated scalar quantities, and regularization of the decomposition.

Let $\mathcal{A} = \{\Delta, S, E, C, \nu\}$ index the annotated scalars—the substantive difference, surface change, and evidence sufficiency read off the auxiliary heads, the reconstructed coverage $\hat{C}_U$, and the continuous novelty score. Each contributes a smooth- L_1 level loss $\mathcal{L}_a$ against its annotation. Level losses alone leave these heads free to settle on the median of their target, so we also supervise their order: the ranking terms of Δ, S, E , and C are collected into a single $\mathcal{L}_{\text{rank}}$, while the novelty score carries its own ranking term inside $\mathcal{L}_\nu$. The objective is

$$\mathcal{L} = \sum_{a \in \mathcal{A}} \lambda_a \mathcal{L}_a + \lambda_{\text{rank}} \mathcal{L}_{\text{rank}} + \lambda_{\text{nov}} \mathcal{L}_{\text{nov}} + \lambda_{\text{type}} \mathcal{L}_{\text{type}} + \lambda_{\text{dec}} \mathcal{L}_{\text{dec}} + \lambda_{\text{var}} \mathcal{L}_{\text{var}}. \quad (7)$$

Here $\mathcal{L}_{\text{nov}}$ supervises the three novelty branches, their fused distribution, and the moderate-to-high head, and $\mathcal{L}_{\text{type}}$ supervises residual categories and nontrivial combinations. Beyond its level loss, the coverage term $\mathcal{L}_C$ adds a consistency loss against the fixed Stage A prediction, assigning less consistency weight to predictions with higher entropy. $\mathcal{L}_{\text{dec}}$ regularizes the decomposition against redundant token assignments and unused units, and $\mathcal{L}_{\text{var}}$ is a variance floor that keeps the unit coverages from coinciding. Appendix C.4 defines each term and reports the coefficients λ .

4 EXPERIMENTS

The evaluation set is the IDEADELTA test split. Every system receives the same evidence — the five highest-ranked admissible priors per target, as titles and abstracts, with no paper text — and is scored against the same labels by the same script, so a difference between two rows is a difference in how the evidence is used. On a four-level label with unequal support, accuracy rewards the majority class and MAE the middle of the scale, while quadratic weighted kappa is what a constant predictor cannot win, so we report all four. Appendix D.1 gives the splits, hyperparameters and scoring code.

4.1 MAIN RESULTS

We compare against a *direct* single call, the *RINoBench* (Schopf & Färber, 2026) and *UKPLab 3-stage* (Afzal et al., 2026) prompt chains, *GraphMind* (da Silva et al., 2025), which labels each prior

Table 2: Baselines and our model on the IdeaDelta evaluation set, 3,352 ideas, same labels and the same idea-only protocol (top-5 retrieved priors, no paper text), split by domain. A row group isolates the effect of the model behind an unchanged prompt chain. A majority-class predictor scores 0.312, 0.356 and 0.328 accuracy with kappa zero by construction, and several rows fall below it; MAE is lowest for rows that concentrate near the middle of the scale, so read it beside QWK.

System	Backbone	Graph ML				Computer vision				Time series			
		Acc	Ma-F1	QWK	MAE ↓	Acc	Ma-F1	QWK	MAE ↓	Acc	Ma-F1	QWK	MAE ↓
Direct, 5 priors	Qwen3.5	0.300	0.266	0.319	0.925	0.308	0.264	0.300	0.876	0.291	0.283	0.327	0.956
	deepseek-flash	0.438	0.444	0.558	0.706	0.409	0.418	0.521	0.728	0.416	0.415	0.560	0.716
	gpt-5.6-luna	0.432	0.439	0.579	0.657	0.420	0.430	0.561	0.661	0.445	0.436	0.576	0.653
RINoBench	Qwen3.5	0.324	0.286	0.335	0.831	0.346	0.293	0.338	0.777	0.344	0.333	0.413	0.793
	deepseek-flash	0.443	0.458	0.560	0.639	0.416	0.411	0.501	0.676	0.442	0.447	0.553	0.654
	gpt-5.6-luna	0.391	0.377	0.487	0.698	0.393	0.364	0.458	0.683	0.392	0.379	0.512	0.688
UKPLab 3-stage	Qwen3.5	0.280	0.161	0.125	0.928	0.289	0.150	0.106	0.873	0.274	0.180	0.132	0.973
	deepseek-flash	0.351	0.267	0.301	0.761	0.376	0.272	0.259	0.741	0.386	0.289	0.290	0.727
	gpt-5.6-luna	0.318	0.226	0.266	0.800	0.367	0.248	0.304	0.721	0.365	0.267	0.296	0.743
GraphMind	Qwen3.5	0.264	0.205	0.154	1.027	0.264	0.182	0.108	0.959	0.244	0.204	0.155	1.051
	deepseek-flash	0.261	0.217	0.154	1.111	—	—	—	—	—	—	—	—
	gpt-5.6-luna	0.366	0.346	0.346	0.778	—	—	—	—	—	—	—	—
DeepReviewer	DeepReviewer-7B	0.294	0.237	0.034	0.965	0.325	0.235	−0.043	0.972	0.296	0.222	−0.017	0.964
EGRD (ours)	—	0.799	0.801	0.909	0.201	0.761	0.767	0.882	0.242	0.789	0.776	0.898	0.215

as supporting or contrasting before rating, and *DeepReviewer-7B* (Zhu et al., 2025), a reviewer model fine-tuned on paper review. The first four appear once per backbone. Table 2 reports every system on the IDEADelta evaluation set, and Table 3 carries the same frozen checkpoint to RINoBench. In the first table the prompted systems occupy a band well below our row.

Among the prompted systems, the backbone matters more than the prompt chain. Holding the prompt fixed and changing the model moves the result further than the reverse. On Graph ML a single five-prior call rises from 0.300 to 0.438 accuracy when Qwen3.5 is replaced by deepseek-flash, a larger move than separates any two prompt chains on one backbone. GraphMind is the exception: on Qwen3.5 it reaches 0.264 against 0.300 for a direct call on the same backbone, at 476 seconds and 31,201 LLM calls for the 3,352 ideas. Its polarity stage is why: given the same five priors Qwen3.5 calls 4.33 supporting against deepseek-flash’s 1.04, and the ratings follow.

The domain, by contrast, does not matter. Within any row the three domains differ by a few points while the gap to EGRD exceeds thirty on every metric, so what separates the two blocks is not a property of one field.

Fine-tuning on paper review does not transfer to ideas. DeepReviewer-7B is the one system trained rather than prompted, and its kappa of 0.034, −0.043 and −0.017 says it does not order the ideas. A model trained to review whole papers appears to read a longer related-work block as evidence of less novelty largely independently of what that work says, which is what scoring a target against a retrieved set read as a block permits.

We evaluate the frozen checkpoint on RINoBench (Schopf & Färber, 2026), whose 277 test ideas carry expert 1–5 labels folded to four levels. These experts never saw the scheme, and the drop between the two tables is the cost of changing annotator.

The ordering transfers even though the level does not. EGRD leads Macro-F1 and kappa at 0.3407 and 0.2245 against 0.2457 and 0.1469 for the best other system on each, and is the only row above 0.15 kappa, while six of the eleven baselines sit below 0.05 and so do not order the ideas at all. Accuracy carries no information here, since a constant predictor scores 0.4152; our row reaches the same 0.4152 with a prediction that is not constant, 16/36/76/149 ideas across the four classes against a gold 15/60/87/115. The margin is modest: we read it as some agreement between residual structure learned on model-generated labels and expert judgment, not as the problem solved.

Two properties of the transfer bound what that agreement can be, and both are visible before any model is run. The first is coverage. IDEADelta is built from three fields, whereas RINoBench draws its ideas from OpenReview submissions across machine learning as a whole, so most of its

Table 3: RINoBench test, 277 ideas, human expert 1–5 labels folded to four levels with 4/5 merged. The merged top class holds 41.5% of the ideas, so predicting it everywhere scores 0.4152 accuracy, at or above every system in the table. Only about 30 of the 277 ideas fall in the three fields EGRD was trained on. Bold marks the best value in a column; Acc and MAE are not led by our row.

System	Acc	Ma-F1	QWK	MAE ↓
<i>Same Qwen3.5 backbone, zero-shot</i>				
Qwen3.5-397B-A17B	0.3574	0.2244	0.1451	0.7726
GraphMind	0.3357	0.2281	0.1469	0.7653
UKPLab 3-stage	0.3466	0.1892	0.0808	0.7220
<i>Released predictions of the benchmark authors, folded to four levels</i>				
GPT-5	0.3646	0.2457	0.0955	0.8051
o3	0.3971	0.2370	0.0606	0.8087
gpt-oss-120b	0.4152	0.2249	0.0392	0.8375
Llama-3.1-8B-Instruct	0.3755	0.1961	−0.0011	0.8773
DeepSeek-R1	0.4224	0.1873	0.0283	0.8700
Llama-4-Scout-17B-16E	0.4116	0.1864	0.0215	0.8809
Llama-3.3-70B-Instruct	0.4188	0.1531	0.0051	0.9025
<i>Reviewer model fine-tuned on paper review, local weights</i>				
DeepReviewer-7B	0.1913	0.1425	0.0145	1.3610
EGRD (ours)	0.4152	0.3407	0.2245	0.8087

Table 4: Ablations on the pooled IDEADELTA evaluation set, 3,352 ideas. **Full** is the unablated checkpoint of Table 2. The pipeline-stage rows intervene on that trained checkpoint at inference and the decomposition-design rows retrain the model without the component, so magnitudes are comparable within a block but not across the two.

System	Acc	Ma-F1	QWK	High-F1
Full	0.7796	0.7790	0.8958	0.7708
<i>Pipeline stages</i>				
Permuting prior evidence	0.3471	0.3532	0.2241	0.4420
Permuting coverage	0.6649	0.6605	0.5996	0.6630
Zeroing residuals	0.6371	0.6554	0.7878	0.7079
<i>Decomposition design</i>				
Collapsing semantic units	0.6580	0.6292	0.8591	0.6120
Disabling slot competition	0.6701	0.6874	0.7665	0.6813
Removing noisy-OR support	0.6210	0.5520	0.6291	0.5665

test ideas sit in areas — robust reinforcement learning, tabular learning, computational biology — whose prior work the model was never trained to read. The second is granularity. The expert scale has five levels against the four EGRD predicts, so the top two are merged, and that merge does two things at once: it puts 41.5% of the ideas into a single class, which is why accuracy is uninformative here, and it erases a distinction the experts themselves chose to draw, so a model that orders the ideas correctly still has no way to express the difference between a 4 and a 5.

The blockwise rule lets the model read more evidence than training supplied. RINoBench offers far more candidate priors than the five EGRD was trained with. Splitting them into disjoint blocks of five and taking the minimum substantive delta over blocks is conservative — any block finding the idea covered lowers its grade — and needs no retraining, since a block has the shape seen in training. Raising the budget from 25 to 60 priors moves accuracy from 0.4116 to 0.4152 and kappa from 0.2182 to 0.2245: real but small, and saturating, since the median idea carries 23 candidates. The same change separates the systems: Qwen3.5 improves from 0.2244 to 0.2477 Macro-F1 while DeepReviewer-7B falls from 0.1425 to 0.1230. A system that treats the amount of retrieved prior work as the signal degrades when given more of it.

4.2 ABLATION

Because the evidence is supplied rather than retrieved, it can be intervened on with the target held fixed, which separates a model that reads the evidence from one that recognises gradable phrasing.

Every stage of the pipeline is load-bearing, and error accumulates along it. The first three interventions cut the chain at successive points. **Permuting prior evidence** shuffles the retrieved prior texts across targets, destroying the semantic content the chain is built on; **permuting coverage** shuffles only the joint coverage score and leaves every prior text in place, destroying what the historical evidence inputs (§3.1) infer from that content; **zeroing residuals** sets z_U , z_I and z_C to zero, destroying what the evidence-grounded residual decomposition (§3.2) adds on top of both. The ablations show that all three are indispensable. Permuting prior evidence costs 0.43 accuracy and 0.67 kappa, dropping accuracy into the band of the majority-class predictor of Table 2 (0.312–0.356): the grade is not recoverable from the target alone, so the priors are read rather than treated as a pretext for grading the target’s own phrasing. Permuting coverage costs 0.11 accuracy and 0.30 kappa with the texts intact, so what the historical evidence inputs infer is not redundant with the evidence it is computed from. Zeroing residuals costs 0.14 accuracy, 0.12 Macro-F1 and 0.11 kappa, so the residual decomposition is not an ornament on a prediction the evidence inputs already support. Taken in order, the damage grows the earlier the cut falls: what one stage corrupts it passes on, and the corruption is amplified rather than absorbed downstream. Accuracy alone reverses at the last step, 0.14 for zeroing residuals against 0.11 for permuting coverage, and it reverses for a reason: the residual decomposition moves the metric that counts exact grades far more than the one that counts agreement in rank, so what it contributes is which grade an idea receives rather than how the ideas are ordered.

The mechanisms inside the residual decomposition earn their place, and two of the three act on the ordering. The remaining interventions leave the pipeline intact and remove one mechanism from the decomposition itself. **Collapsing semantic units**, from eight to one, costs 0.12 accuracy, 0.15 Macro-F1 and 0.16 High-F1 but only 0.04 kappa, so the collapse loses the grade and leaves the ordering nearly intact. **Disabling slot competition**, which removes the softmax through which the units compete for each token, costs 0.11 accuracy and 0.13 kappa, reversing that pattern: deciding which unit a token belongs to is what keeps the units apart in the first place. **Removing noisy-OR support** costs the most of the three on every metric, 0.16 accuracy, 0.23 Macro-F1, 0.27 kappa and 0.20 High-F1, so aggregating support across priors is what makes the ordering work. For the collapsed units and the removed noisy-OR, High-F1 falls further than accuracy, which is where a decomposition should matter: the top grade turns on parts of a target that no single prior supports.

We also sweep the three parameters that govern how much the model is given (the evidence budget K , the training fraction and the encoder input length) one at a time from the default configuration of Table 2, in all three domains. The largest of the three moves accuracy by at most 0.075, against the 0.34 that separates our row from the best prompted row of Table 2, so none of them is what the main comparison turns on. The sweeps, and what they do and do not license, are in Appendix D.2.

REFERENCES

- Osama Mohammed Afzal, Preslav Nakov, Tom Hope, and Iryna Gurevych. Beyond “not novel enough”: Enriching scholarly critique with LLM-assisted feedback. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2648–2671, Rabat, Morocco, 2026. Association for Computational Linguistics. URL <https://aclanthology.org/2026.eacl-long.121/>.
- Jinheon Baek, Sujay Kumar Jauhar, Silviu Cucerzan, and Sung Ju Hwang. ResearchAgent: Iterative research idea generation over scientific literature with large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 6709–6738, Albuquerque, New Mexico, 2025. Association for Computational Linguistics. URL <https://aclanthology.org/2025.naacl-long.342/>.
- Arman Cohan, Sergey Feldman, Iz Beltagy, Doug Downey, and Daniel Weld. SPECTER: Document-level representation learning using citation-informed transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 2270–2282,

- Online, 2020. Association for Computational Linguistics. URL <https://aclanthology.org/2020.acl-main.207/>.
- Italo Luis da Silva, Hanqi Yan, Lin Gui, and Yulan He. GraphMind: Interactive novelty assessment system for accelerating scientific discovery. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, 2025. URL <https://arxiv.org/abs/2510.15706>.
- Maximilian Ilse, Jakub M. Tomczak, and Max Welling. Attention-based deep multiple instance learning. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pp. 2127–2136, 2018. URL <https://proceedings.mlr.press/v80/ilse18a.html>.
- Weixin Liang, Yuhui Zhang, Hancheng Cao, Binglu Wang, Daisy Yi Ding, Xinyu Yang, Kailas Vodrahalli, Siyu He, Daniel Scott Smith, Yian Yin, Daniel McFarland, and James Zou. Can large language models provide useful feedback on research papers? a large-scale empirical analysis. *NEJM AI*, 1(8), 2024. doi: 10.1056/AIoa2400196.
- Francesco Locatello, Dirk Weissenborn, Thomas Unterthiner, Aravindh Mahendran, Georg Heigold, Jakob Uszkoreit, Alexey Dosovitskiy, and Thomas Kipf. Object-centric learning with slot attention. In *Advances in Neural Information Processing Systems* 33, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/8511df98c02ab60aealb2356c013bc0f-Abstract.html>.
- Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024. URL <https://arxiv.org/abs/2408.06292>.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 12076–12100, Singapore, 2023. Association for Computational Linguistics. URL <https://aclanthology.org/2023.emnlp-main.741/>.
- Malte Ostendorff, Nils Rethmeier, Isabelle Augenstein, Bela Gipp, and Georg Rehm. Neighborhood contrastive learning for scientific document representations with citation embeddings. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 11670–11688, Abu Dhabi, United Arab Emirates, 2022. Association for Computational Linguistics. URL <https://aclanthology.org/2022.emnlp-main.802/>.
- Tim Schopf and Michael Färber. Is this idea novel? an automated benchmark for judgment of research ideas. In *Proceedings of the 2026 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC)*, 2026. URL <https://arxiv.org/abs/2603.10303>.
- Chenglei Si, Diyi Yang, and Tatsunori Hashimoto. Can LLMs generate novel research ideas? a large-scale human study with 100+ NLP researchers. *arXiv preprint arXiv:2409.04109*, 2024. URL <https://arxiv.org/abs/2409.04109>.
- Brian Uzzi, Satyam Mukherjee, Michael Stringer, and Ben Jones. Atypical combinations and scientific impact. *Science*, 342(6157):468–472, 2013. doi: 10.1126/science.1240474.
- David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 7534–7550, Online, 2020. Association for Computational Linguistics. URL <https://aclanthology.org/2020.emnlp-main.609/>.
- David Wadden, Kyle Lo, Lucy Lu Wang, Arman Cohan, Iz Beltagy, and Hannaneh Hajishirzi. MultiVerS: Improving scientific claim verification with weak supervision and full-document context. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pp. 61–76, Seattle, United States, 2022. Association for Computational Linguistics. URL <https://aclanthology.org/2022.findings-naacl.6/>.

- Qingyun Wang, Doug Downey, Heng Ji, and Tom Hope. SciMON: Scientific inspiration machines optimized for novelty. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 279–299, Bangkok, Thailand, 2024. Association for Computational Linguistics. URL <https://aclanthology.org/2024.acl-long.18/>.
- Miriam Wanner, Seth Ebner, Zhengping Jiang, Mark Dredze, and Benjamin Van Durme. A closer look at claim decomposition. In *Proceedings of the 13th Joint Conference on Lexical and Computational Semantics (*SEM 2024)*, pp. 153–175, Mexico City, Mexico, 2024. Association for Computational Linguistics. URL <https://aclanthology.org/2024.starsem-1.13/>.
- Jerry Wei, Chengrun Yang, Xinying Song, Yifeng Lu, Nathan Hu, Jie Huang, Dustin Tran, Daiyi Peng, Ruiibo Liu, Da Huang, Cosmo Du, and Quoc V. Le. Long-form factuality in large language models. In *Advances in Neural Information Processing Systems 37*, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/937ae0e83eb08d2cb8627feldef8c751-Abstract-Conference.html.
- Minjun Zhu, Yixuan Weng, Linyi Yang, and Yue Zhang. DeepReview: Improving LLM-based paper review with human-like deep thinking process. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 29330–29355, Vienna, Austria, 2025. Association for Computational Linguistics. URL <https://aclanthology.org/2025.acl-long.1420/>.

A RELATED WORK

Automated research and idea generation. Language models are increasingly used to carry out parts of the research process itself. Wang et al. (2024) retrieve inspirations from past literature and optimise explicitly for novelty by comparing a candidate against prior papers; Baek et al. (2025) refine proposals against a set of LLM reviewing agents that give structured feedback along several dimensions; and Lu et al. (2024) close the loop from ideation through experiments to a written paper and a simulated review. Si et al. (2024) evaluate what such pipelines produce, finding the generated ideas rated more novel than human-written ones but less feasible. Every one of these systems contains an internal novelty check, and it is the check rather than the generator that bounds them: once proposals are cheap, the output of the pipeline is whatever its filter admits. The filter is nonetheless the component these systems evaluate least directly, judged only through the ideas that eventually emerge downstream of it.

Judging the novelty of scientific work. Studied on its own, the check has passed through several generations of method. The earliest operate on citation statistics, measuring how unusual the combinations of references in a paper are against historical co-occurrence (Uzzi et al., 2013); they require the work to already exist and be cited, and reduce novelty to patterns of co-occurrence rather than to what a paper says. Document-level semantic representations replace surface statistics with learned relatedness, scoring how close a submission is to prior papers through citation-informed embeddings (Cohan et al., 2020; Ostendorff et al., 2022), but they compare documents as points and recover nothing about which content is shared. Most recent systems retrieve related literature and hand the comparison to a language model. Liang et al. (2024) obtain structured feedback from a single prompt, of which significance and novelty is one section, and report partial overlap with human reviews; Zhu et al. (2025) train a reviewer model on a multi-stage process of structured analysis, retrieval, and evidence-based argumentation; Afzal et al. (2026) split the judgment into claim extraction, retrieval and synthesis of related work, and structured comparison, motivating the design with an analysis of human-written novelty reviews; and da Silva et al. (2025) build an interactive system on the same retrieve-then-compare pattern. Schopf & Färber (2026) supply the evaluation counterpart and report that model ratings still diverge substantially from expert ones even when the accompanying reasoning appears sound. Two properties persist across these generations, and both are structural rather than incidental. The evidence is retrieved by the system under evaluation, so its verdicts are unattributable: an error cannot be assigned to the retrieval or to the judgment, and the two therefore cannot be improved, or even measured, separately. We remove this by conditioning on a given, temporally admissible evidence set, which additionally turns the evidence into a variable that can be intervened on. The second property is that the target is scored as an undivided whole,

against the closest prior work or against a retrieved set read as a block. Such a comparison returns one number per target, and no post-processing of that number can recover which part of the idea it is about; the information is discarded before the score is formed, not lost afterwards.

Decomposed verification against evidence. Saying which part remains requires decomposing the target, and outside novelty assessment this has become the standard recipe for factuality evaluation: break a text into smaller propositions and check each against retrieved evidence. Min et al. (2023) decompose a passage into atomic facts, verify each against a knowledge source, and report the supported fraction; Wei et al. (2024) scale the same procedure with search as the evidence source; and in the scientific domain the corresponding task checks a stated claim against candidate abstracts while identifying the supporting rationales (Wadden et al., 2020; 2022). Three properties of the recipe do not survive the move to research ideas. The propositions are assumed to be given, whereas an idea carries no stated claims, and the decomposition step is not a neutral preprocessor: Wanner et al. (2024) show that scores of this family shift materially with the decomposition used. Support is treated as a label, whereas a prior work can cover part of a unit and not the rest. And each unit is verified independently before the results are averaged, so the score is a function of the per-unit coverages alone and cannot express whether a combination of units has precedent. We accordingly induce the units rather than reading them off the text, measure support as a graded quantity, and add a residual defined over pairs of units.

B DATASET DETAILS

B.1 DATASET CONSTRUCTION PIPELINE

Sources and idea extraction. Papers are drawn from arXiv in three areas—graph machine learning, computer vision, and time series. Selection prefers work that was subsequently published and, among that, papers at major venues. From each paper’s title, abstract and introduction an LLM extracts innovation claims and assigns each an innovation aspect; a second LLM pass discards claims that are not innovation claims at all—open-source and reproducibility statements, bare state-of-the-art assertions, and engineering deliverables—so that later stages judge substantive claims rather than boilerplate.

Candidate priors. Idea texts are embedded and indexed for top- K nearest-neighbour recall, and candidate relations are formed under a hybrid rule: high similarity is accepted directly, low similarity is rejected directly, and the uncertain band is passed to an LLM. The base build pairs ideas of the same type only, which would record an idea whose closest relative is of another type—a theoretical contribution nearest to a method—as isolated; a cross-type pass recovers these. At most $K = 5$ ranked priors are retained per target. Pairing stays within a source domain, and no cross-domain pair is labelled, so the three single-domain subsets and their union address the same ideas under identifiers namespaced by domain. The assembled dataset passes its own validation checks with no duplicate pair keys, no duplicate novelty targets, no invalid pair labels, no target left without a pair, and no temporal violation.

Annotation. All annotation stages use the same judge, `deepseek-v4-pro`. It labels pair coverage over the four classes, with the ambiguous band—related but not covering, partial cover, large cover, and high-similarity pairs the graph left unconnected—re-judged rather than accepted from the similarity rule. It then estimates, for each target, the substantive residual that survives the retained evidence, the sufficiency of that evidence, and the extent to which the claimed difference is surface-level. Idea-level grades follow by applying Eq. (1) to those quantities. Labels are model-generated; each stage was reviewed by sampled manual inspection.

Splits and balancing. Splits are date-ordered rather than random, so that no training target is dated after an evaluation target. The cuts fall at 2025-02 and 2025-11, giving 59,269/16,436/17,171 pairs and 13,815/3,944/3,352 ideas. The cut months are shared at the boundary, so the partition is ordered but not strictly month-disjoint, and isolation at the level of source papers has not been established. A balanced training variant resamples the pair training set to 1,000 examples per coverage class, at most five copies of any example and marked as duplicates; dev and test keep their natural distribution

and are never resampled. Pair coverage labels are skewed in the natural distribution: of 92,876 pairs, 41,765 are not covering, 38,775 related but not covering, 9,582 partial cover, and 2,754 large cover.

B.2 VALIDATION OF RESEARCH NOVELTY LABELS

Our novelty labels assess the substantive contribution of a target idea beyond its retrieved prior work. We validate this interpretation through complementary experiments. First, we examine how well metadata and retrieved prior content predict the existing labels. We then modify the evidence and repeat annotation to test whether novelty decreases when the target contribution is already covered and remains stable when its substantive contribution is unchanged.

Prediction from metadata. We first test whether simple metadata can account for the novelty labels without describing the technical content of prior work. We consider seven variables: character length, word count, publication year, three research-area attributes, and the number of retrieved priors. Over all 21,111 ideas, we measure their individual associations with the novelty grade using Spearman ρ for continuous variables and ε^2 for categorical variables. We set the thresholds in advance at $|\rho| \geq 0.20$ and $\varepsilon^2 \geq 0.06$.

Text length, word count, and publication year show weak associations ($\rho = +0.049$, $+0.058$, and -0.130 , respectively), while the three research-area attributes have $\varepsilon^2 \leq 0.033$. All remain below the corresponding thresholds. The number of priors and their mean coverage are negatively correlated with novelty ($\rho = -0.286$ and -0.471 , respectively). The negative association with coverage is consistent with the intended meaning of the labels: contributions already covered by prior work receive lower novelty grades.

Individual associations may overlook information available from several variables together. We therefore train a classifier on all seven metadata variables to predict the four novelty grades. We use the date-ordered split in Section 2.3, with 13,815 training ideas and 3,352 test ideas, and evaluate quadratic weighted kappa (QWK), which accounts for the ordering of the grades. For comparison, we train the same classifier using only mean prior coverage. QWK increases from 0.227 with metadata to 0.539 with coverage alone. Coverage thus predicts the labels more accurately than the metadata considered here, supporting the relevance of the relationship between the target and prior work.

Prediction with retrieved and random priors. We next examine whether the specific retrieved content matters for label prediction. In the TRUE condition, each target is paired with its five retrieved priors. In the RANDOM condition, these are replaced by five ideas sampled uniformly from those published before the target. The two conditions use the same number of priors and total evidence text volume, and all priors precede the target. We hold the model, hyperparameters, date-ordered split, and random seeds fixed. Models are trained separately from scratch and evaluated on the same held-out targets against their original labels.

Across the three seeds, QWK decreases from 0.196 with retrieved priors to 0.005 ± 0.021 with random priors. The gap of 0.192 is about an order of magnitude larger than the between-seed standard deviation of 0.015. Replacing retrieved priors with random earlier work therefore reduces prediction to near-chance agreement. This comparison supports the importance of the retrieved content, beyond the amount of evidence supplied.

Response to evidence covering the target contribution. The prediction experiments keep the labels fixed. We now test whether the annotation procedure responds appropriately when the evidence changes. If an earlier study already contains the target contribution, the assigned novelty should decrease. We sample 30 targets across the four grades and construct one synthetic prior for each target. This prior describes the same technical contribution and is dated before the target. After adding it to the evidence set, we repeat the complete annotation procedure: pair coverage labeling, residual contribution estimation, and grading according to Eq. (1). Synthetic priors have dedicated identifiers and are marked `synthetic_control`. They do not refer to real papers, authors, or DOIs and are excluded from the released dataset.

Each modified target is compared with five independent annotations of its original evidence set to account for variation across repeated annotations. We evaluate recognition of the added prior as

large coverage, changes in joint coverage and residual contribution, and decreases in the novelty score and grade. Ties are excluded from the exact sign test.

The added prior is labeled large coverage in all 29 comparable cases. Joint coverage increases, while residual contribution and novelty score decrease. The novelty grade decreases for 82.8% of targets (sign test $p = 2.2 \times 10^{-7}$). These results show that the annotation procedure lowers novelty when the evidence covers the target contribution.

Stability under changes that preserve the contribution. A novelty judgment should also remain stable when the evidence changes without altering the substantive contribution of the target. For the same 30 targets, we reorder the priors, duplicate existing priors, or add unrelated papers. The novelty grades before and after these changes show high agreement: $\kappa = 0.970, 0.902$, and 1.000 , respectively. These results complement the coverage experiment: the judgments respond to evidence that reduces the remaining contribution and remain comparatively stable under changes that preserve it.

Evaluation scope. The metadata analysis covers 21,111 ideas, and the prediction experiments use 13,815/3,944/3,352 training, validation, and test examples ordered by date. The evidence modification experiments cover 30 targets, with 150 baseline repetitions, 30 coverage comparisons, and 160 stability comparisons. Sampling uses a fixed seed and groups ideas by source paper. Each modified example is compared with its own original version, and confidence intervals for component-level changes use 2,000 bootstrap samples clustered by target.

C METHOD DETAILS

C.1 HISTORICAL COVERAGE MODELS

Pair coverage and evidence selection. For each target–prior pair (x, p_k) , the pair coverage model estimates how much of the target the prior accounts for. It outputs a score $s_k \in [0, 1]$ and a distribution over four classes: unrelated, related but uncovered, partially covered, and largely covered. A cross-encoder predicts relatedness a_k , coverage conditional on relatedness b_k , and large coverage conditional on coverage d_k . These probabilities define the class distribution

$$\pi_k^{\text{pair}} = [1 - a_k, a_k(1 - b_k), a_k b_k(1 - d_k), a_k b_k d_k]. \quad (8)$$

With classes indexed from 0 to 3, the coverage score is $s_k = \frac{1}{2}\pi_{k,2}^{\text{pair}} + \pi_{k,3}^{\text{pair}}$. The reference pipeline ranks temporally admissible candidates by $s_k + 0.02\ell_k/3$, where ℓ_k is the predicted ordinal coverage label, and retains at most K priors. The reported IdeaDelta checkpoint instead uses the ordering and pair annotations supplied in the dataset ranking records, as detailed under the evidence-source protocol below. In both cases, the retained priors and their pair coverage features form the evidence inputs to the subsequent stages.

Pair feature vector. The features carried into Stages B and C are

$$\mathbf{f}_k = [s_k; s_k + 0.02\ell_k/3; \pi_k^{\text{pair}}; a_k; a_k b_k; d_k; \kappa_k/K; \mathbb{I}[\text{prior text available}]] \in \mathbb{R}^{11}, \quad (9)$$

where κ_k is the prior’s rank within $\mathcal{E}_x$. The first seven entries are deterministic functions of (a_k, b_k, d_k) through Eq. (8), so $\mathbf{f}_k$ carries three degrees of freedom from the pair model alongside the two ranking quantities and the availability flag. We keep the redundant parameterization because the downstream projections consume these quantities at different scales.

Global context features. The twenty-one scalars projected into $\mathbf{z}_G$ are the scaled target timestamp; ten statistics of the pair scores over $\mathcal{E}_x$ —the fraction of the K evidence slots that are filled, the maximum, mean, standard deviation and sum of the s_k , a union score over them, the fractions of retained priors judged covering and judged related, the margin between the top two ranking scores, and the entropy of the pair scores; and ten joint-coverage features—two scalar coverage estimates, the raw and normalized entropies of π^J , an indicator for whether a joint prediction is available, the gap between the joint estimate and the pair-score union, and π^J itself.

Joint coverage. Different priors may cover different parts of the target. We therefore estimate their joint coverage with a multiple-instance model, treating each pair (x, p_k) as an instance and the retained evidence set as a bag (Ilse et al., 2018). Let $n = |\mathcal{E}_x|$ denote the number of retained priors. A shared cross-encoder maps each pair to a representation $\mathbf{h}_k^J$, from which the model computes attention weights β_k and internal evidence scores e_k :

$$\beta_k = \text{softmax}_k(\mathbf{w}_a^\top \mathbf{h}_k^J + b_a), \quad e_k = \sigma(\mathbf{w}_e^\top \mathbf{h}_k^J + b_e). \quad (10)$$

For a nonempty evidence set, noisy-OR pools the evidence scores as

$$e_U = 1 - \prod_{k=1}^n (1 - e_k). \quad (11)$$

The model summarizes these scores by their noisy-OR aggregate, maximum, mean, and the fraction of the K evidence slots that are filled. It concatenates these statistics with attention-pooled and elementwise maximum-pooled pair representations:

$$\mathbf{v}_E = \left[e_U, \max_k e_k, \frac{1}{n} \sum_{k=1}^n e_k, \frac{n}{K} \right], \quad (12)$$

$$\mathbf{z}_J = f_J \left(\left[\sum_{k=1}^n \beta_k \mathbf{h}_k^J; \max_k \mathbf{h}_k^J; \mathbf{v}_E \right] \right). \quad (13)$$

Here f_J is an MLP. Attention pooling aggregates evidence across priors, while maximum pooling retains the strongest feature values. A hierarchical head maps $\mathbf{z}_J$ to three probabilities: any coverage, substantial coverage conditional on any coverage, and mostly covered conditional on substantial coverage. Using the factorization in Eq. (8), these probabilities yield π^J over four joint-coverage classes: *not*, *weakly*, *partially*, and *mostly* covered. The scalar joint-coverage estimate is $C_J = \sum_{\ell=0}^3 \mu_\ell \pi_\ell^J$, where μ_ℓ is the mean annotated coverage score for class ℓ on the training split. Both C_J and π^J are supplied as fixed inputs to the residual and novelty modules. The noisy-OR aggregate e_U contributes one of the four statistics in $\mathbf{v}_E$; the joint-coverage prediction is learned from the full bag representation.

C.2 RESIDUAL REPRESENTATIONS AND POOLING

All representations introduced in Stages B and C—the queries $\mathbf{q}_m$, the token states $\mathbf{h}_j$, the units $\mathbf{u}_m$, the pair representations $\mathbf{v}_k$ and $\tilde{\mathbf{v}}_k$, the interaction representations $\mathbf{t}_{ij}$, and the pooled vectors $\mathbf{z}_U$, $\mathbf{z}_I$, $\mathbf{z}_C$, $\mathbf{z}_R$ and $\tilde{\mathbf{h}}_x$ —lie in $\mathbb{R}^d$. Eq. (3) accordingly gives f_A an input of dimension $4d + 11$, f_I takes $4d$, f_R takes $3d$, and the shared representation $\mathbf{h}$ has dimension $4d$.

Pair encoding. For each target–prior pair (x, p_k) , the residual cross-encoder produces token states $\mathbf{H}_k \in \mathbb{R}^{L \times d}$, where L is the input sequence length. We apply masked mean pooling to obtain the pair representation and augment it with projected pair coverage features:

$$\mathbf{v}_k = \text{MeanPool}(\mathbf{H}_k), \quad \tilde{\mathbf{v}}_k = \mathbf{v}_k + f_P(\mathbf{f}_k), \quad (14)$$

where f_P consists of a linear map, layer normalization, GELU, and dropout. The residual encoder shares parameters across all retained priors. To construct the target token bank, we average target-token states at matching positions across valid pairs. Masked mean pooling over this bank yields the target representation $\tilde{\mathbf{h}}_x$ used in residual fusion.

Coverage reconstruction. To supervise the unit coverage estimates, we aggregate them into a scalar joint-coverage prediction. The pooling weights depend on both the semantic representation and the activity of each unit:

$$\hat{C}_U = \sum_{m=1}^M \omega_m c_m, \quad \omega_m = \frac{v_m \exp(\mathbf{w}_c^\top \mathbf{u}_m + b_c)}{\sum_{r=1}^M v_r \exp(\mathbf{w}_c^\top \mathbf{u}_r + b_c)}. \quad (15)$$

The resulting estimate $\hat{C}_U$ is supervised by annotated coverage and the fixed Stage A prediction. This aggregate supervision constrains unit coverage but does not uniquely determine the latent support matrix. We use noisy-OR as a differentiable rule for aggregating support; its use does not establish that support from different priors is statistically independent.

Interaction encoding and attention. For each unit pair $i < j$, an MLP f_I encodes the two units and their elementwise comparisons into an interaction representation. A sigmoid head then predicts the interaction strength:

$$\mathbf{t}_{ij} = f_I([\mathbf{u}_i; \mathbf{u}_j; \mathbf{u}_i \odot \mathbf{u}_j; |\mathbf{u}_i - \mathbf{u}_j|]), \quad g_{ij} = \sigma(\mathbf{w}_g^\top \mathbf{t}_{ij} + b_g). \quad (16)$$

Using the interaction residual weights r_{ij} defined in Stage B, we compute attention over all $\binom{M}{2}$ unit pairs:

$$\eta_{ij} = \text{softmax}_{i < j} (\mathbf{w}_I^\top \mathbf{t}_{ij} + b_I + \log(r_{ij} + \epsilon)), \quad \mathbf{z}_I = \sum_{i < j} \eta_{ij} r_{ij} \mathbf{t}_{ij}. \quad (17)$$

The logarithmic bias directs attention toward pairs with larger residual weights. Multiplication by r_{ij} during pooling also preserves the scale of the residual weights, allowing weak residuals to contribute less even after attention normalization.

Padded evidence slots carry zero support: the support matrix is masked by the evidence-availability flag and clamped to $[0, 0.995]$ before use. Both aggregations are therefore well defined when fewer than K priors are retained—the product in Eq. (4) runs over the retained priors, and the maximum in Eq. (5) is taken over all K slots, so $b_{ij} = 0$ when none is retained.

C.3 RESIDUAL FUSION, PREDICTION HEADS, AND DECISION RULES

Context gate. The fusion gate conditions the contribution of global context on the two residual representations and the uncertainty of the joint-coverage prediction. Let H_J denote the normalized entropy of π^J . The gate is

$$\gamma = \sigma(f_\gamma([\mathbf{z}_U; \mathbf{z}_I; \mathbf{z}_G; H_J])). \quad (18)$$

Here f_γ is a two-layer MLP with a $d/2$ bottleneck and a scalar output, so γ scales $\mathbf{z}_G$ uniformly rather than gating it per dimension. The gated context $\gamma \mathbf{z}_G$ enters residual fusion alongside $\mathbf{z}_U$ and $\mathbf{z}_I$. Global context $\mathbf{z}_G$ also enters the final shared representation directly. Thus, γ controls the context contribution within the residual branch, while the direct context path remains available to the prediction heads.

Auxiliary predictions. The auxiliary heads use the shared representation $\mathbf{h}$ to predict substantive difference $\hat{d}$, surface change $\hat{s}$, evidence sufficiency $\hat{e}$, and residual-type probabilities π^T . The first three outputs are sigmoid scores; the type prediction uses the hierarchy described below.

Residual types. The type hierarchy distinguishes a minor variant from six substantive families: conceptual innovation, mechanism or system, task or application, training or evaluation, data-related innovation, and combination. Conditional fine heads yield eleven types in total. The minor/substantive gate additionally reads $\hat{d}$, $\hat{s}$, and $\max_{i < j} r_{ij}$. The probability assigned to the combination type is multiplied by $0.70 + 0.30\sigma(f_{\text{comb}}(\mathbf{z}_I))$ before the hierarchical distribution is renormalized. The final type distribution is a convex mixture of the hierarchical and flat predictions, with flat weight 0.15.

Novelty branches. All novelty branches receive the prediction features

$$\boldsymbol{\xi} = [\mathbf{h}; \hat{d}; \hat{s}; \hat{e}; \pi^T; C_J; \pi^J]. \quad (19)$$

We encode LOW, WEAK, MODERATE, and HIGH as 0, 1, 2, and 3, respectively. A direct linear classifier and a two-layer MLP classifier map $\boldsymbol{\xi}$ to class distributions $\mathbf{p}^D$ and $\mathbf{p}^R$.

The ordinal branch models cumulative probabilities of reaching each grade threshold. It produces three sigmoid gates g_1, g_2, g_3 from $\boldsymbol{\xi}$ and defines

$$q_1 = g_1, \quad q_2 = g_1 g_2, \quad q_3 = g_1 g_2 g_3, \quad \mathbf{p}^O = [1 - q_1, q_1 - q_2, q_2 - q_3, q_3]. \quad (20)$$

Here q_ℓ estimates the probability that $y \geq \ell$, with $q_1 \geq q_2 \geq q_3$ by construction.

Probability fusion. We aggregate the three branch distributions through a convex mixture:

$$\mathbf{p} = \lambda_D \mathbf{p}^D + \lambda_O \mathbf{p}^O + \lambda_R \mathbf{p}^R, \quad \lambda_D + \lambda_O + \lambda_R = 1, \quad \lambda_D, \lambda_O, \lambda_R \geq 0. \quad (21)$$

The fused distribution provides the initial grade for the refinement below.

Moderate-to-high refinement. A separate binary head $h_{MH} = \sigma(f_{MH}(\xi))$ scores the distinction between moderate and high novelty. At inference, let $y_0 = \arg \max_{\ell} p_{\ell}$ for the mixture in Eq. (21), and define

$$h_{\text{high}} = \alpha_D p_3^D + \alpha_O p_3^O + \alpha_{MH} h_{MH}, \quad \hat{y} = \begin{cases} 3, & y_0 = 2 \text{ and } h_{\text{high}} \geq \theta_{MH}, \\ y_0, & \text{otherwise.} \end{cases} \quad (22)$$

where $\alpha_D, \alpha_O, \alpha_{MH} \geq 0$ and $\alpha_D + \alpha_O + \alpha_{MH} = 1$. The refinement applies a threshold to this weighted score and does not require separate agreement from all three heads. For example, an initial MODERATE prediction is promoted to HIGH if the combined score reaches θ_{MH} , even if the individual heads favor different grades. The distribution-mixture weights and refinement threshold are selected on development data and fixed for testing. Appendix C.4 reports the implementation values.

Continuous outputs. A separate sigmoid head predicts a continuous novelty score. At inference, novelty grades are determined by probability fusion and the moderate-to-high refinement; the continuous score output does not enter that decision rule. The continuous score head and the substantive-difference head are separate readouts; the latter is used in the blockwise RINoBench evaluation.

C.4 TRAINING AND IMPLEMENTATION DETAILS

This section supplements the training objective in Eq. (7) with the loss definitions and implementation settings for the prediction modules of Appendix C.3 and the residual decomposition of Appendix C.2.

Supervision. The classification group $\mathcal{L}_{\text{nov}}$ trains the direct and MLP classifiers with class-weighted cross-entropy, the ordinal outputs against $\mathbf{1}[y \geq \ell]$, $\ell = 1, 2, 3$, with weighted binary cross-entropy, and the fused distribution with a negative log-likelihood. The moderate-to-high head sees only moderate and high examples and adds a positive-negative margin to its binary cross-entropy, weighting moderate examples scored as high more heavily. The type group $\mathcal{L}_{\text{type}}$ combines the hierarchy’s conditional losses, a flat classification loss, binary supervision for nontrivial combinations, and consistency penalties tying the combination branch to the type predictions. The separate novelty score regresses ν^* , the continuous annotation target defined by $\text{raw}(x)$ in Eq. (1) and formed from the annotated coverage, substantive difference, evidence sufficiency, and surface change of the target. It is the quantity the dataset thresholds into the four grades.

Scalar heads: level and ranking terms. Let ℓ_{δ} denote smooth- L_1 loss and w_x the example weight. For $a \in \mathcal{A}$, the level loss is $\mathcal{L}_a = \sum_x w_x \ell_{\delta}(\hat{a}_x - a_x^*) / \sum_x w_x$, where a_x^* is the annotation and $\hat{a}_x$ the prediction. The corresponding ranking term takes within-minibatch pairs separated by $|a_x^* - a_{x'}^*| \geq 0.05$ and contributes $[0.02 - \text{sgn}(a_x^* - a_{x'}^*)(\hat{a}_x - \hat{a}_{x'})]_+$, weighted by $\sqrt{w_x w_{x'}}$; it is dropped when a minibatch holds no eligible pair. The ranking terms of Δ , S , E , and C are summed into $\mathcal{L}_{\text{rank}}$ with the internal weights of Table 5; the novelty score’s ranking term is carried inside $\mathcal{L}_{\nu}$ instead, where it takes most of the weight because the grade is read off the score by thresholding and therefore depends only on order.

The ranking terms are not redundant with the level losses. At the transition δ and target spread used here the level losses operate in their L_1 regime, whose optimum is the median of the target, and a head supervised by the level loss alone drifts towards emitting a near-constant value; the evidence head is the clearest case. Ranking is scale-free and counteracts this.

Coverage term. Beyond its level loss against the annotated joint value C^* , $\mathcal{L}_C$ adds a consistency loss against the upstream score C_J at relative weight 0.2, reweighted by $1 - H_J$ so that examples with uncertain upstream coverage count for less. No observed unit-level support matrix is supplied as a target.

Decomposition regularization and variance floor. $\mathcal{L}_{\text{dec}}$ combines three terms. Attention regularization penalizes the squared entries of $\tilde{\alpha}\tilde{\alpha}^{\top} - I$ for the row-normalized attention matrix $\tilde{\alpha}$, plus half the mean off-diagonal cosine similarity. Query diversity penalizes the mean squared cosine

Table 5: Residual-stage settings for the reported IdeaDelta checkpoint. The checkpoint was selected on development data at epoch 5 of 8.

Setting	Value
Training batch size / evaluation batch size	2 / 8
Head learning rate / encoder learning rate	$1.5 \times 10^{-5} / 3 \times 10^{-6}$
Weight decay / dropout	0.01 / 0.15
Smooth- L_1 transition δ	0.25
$\lambda_\Delta / \lambda_S / \lambda_E$	1.0 / 0.6 / 0.8
λ_C / λ_ν	0.3 / 0.6
$\lambda_{\text{nov}} / \lambda_{\text{type}}$	0.8 / 1.0
λ_{rank}	0.5
$\lambda_{\text{dec}} / \lambda_{\text{var}}$	0.05 / 0.02
Within $\mathcal{L}_{\text{rank}}$: $\Delta / S / E / C$	1.0 / 0.5 / 0.5 / 0.5
Within $\mathcal{L}_\nu$: ranking weight	2.0
Within $\mathcal{L}_C$: Stage A consistency weight	0.2
Within $\mathcal{L}_{\text{dec}}$: attention / query / usage	0.7 / 0.1 / 0.2
Variance floor ν	0.0025
Seed	173

similarity between distinct queries. A usage penalty $D_{\text{KL}}(\rho \| \mathbf{u})$, with $\rho_m = \sum_j B_{mj} / \sum_{r,j} B_{rj}$ and $\mathbf{u}$ uniform over the M units, keeps units from going unused. Their relative weights are given in Table 5.

The variance floor is a separate term rather than a component of $\mathcal{L}_{\text{dec}}$, and enters Eq. (7) under its own coefficient:

$$\mathcal{L}_{\text{var}} = \mathbb{E}_x [\nu - \text{Var}_m(c_{x,m})]_+ / \nu. \quad (23)$$

It acts on the unit coverages rather than on the token assignments, and keeps the c_m within an example from coinciding. The reported configuration uses softmax competition with dot-product scores.

Curriculum and feature perturbation. For a randomly sampled fraction of training examples the joint-coverage features, together with the associated probability and entropy features, are replaced by values derived from annotated joint coverage, with probability 0.50 in epoch 1, 0.20 in epoch 2, and 0 thereafter. This is training-time supervision only: development and test runs always use the upstream joint predictions. Selected joint-coverage features on the global-context path also receive Gaussian noise of standard deviation 0.03 and independent feature dropout at probability 0.10, which leaves untouched both the direct coverage inputs to the novelty heads and the entropy input to the fusion gate.

Inputs and initialization. Pair probabilities missing from a supplied ranking record are reconstructed by the input adapter from the available coverage label and score; pair labels and scores also appear in the pair text input. The two facet embeddings entering $\mathbf{z}_G$ have dimension 32 each. The model uses DeBERTa-v3-base with hidden dimension 768, maximum pair-input length 384, $K = 5$, and $M = 8$; the residual encoder is initialized from the learned joint model and upstream checkpoints stay fixed. Padded priors have zero support, support values are capped at 0.995 for numerical stability, and an empty evidence set falls back to a target-containing input with all prior support masked out. Fusion uses refinement weights $(\alpha_D, \alpha_O, \alpha_{MH}) = (0.40, 0.30, 0.30)$ and mixture weights $(\lambda_D, \lambda_O, \lambda_R) = (0.50, 0.35, 0.15)$ in training; development-set calibration selects $(0.60, 0.20, 0.20)$ and $\theta_{MH} = 0.32$ for inference, fixed thereafter. Table 5 lists the remaining settings.

Evidence-source protocol. The reference pair pipeline predicts coverage features and ranks priors by the rule in Appendix C.1, while the upstream joint model is trained with bag-level coverage supervision and an auxiliary pair-score regression objective. For the reported checkpoint the evidence order and pair annotations come from the dataset ranking records, so the configuration evaluates

novelty prediction with supplied pair annotations and predicted joint coverage. Applying the model to an unseen idea additionally requires computing the pair evidence features.

D EXPERIMENTAL DETAILS

D.1 SETUP

Splits. The split is a date-ordered 13,815/3,944/3,352 partition of the 21,111 idea-level samples in IDEADELTA.

Configuration. All runs share one configuration and vary a single parameter at a time: deberta-v3-base backbone, 8 epochs, batch size 2, learning rate 1.5×10^{-5} with an encoder learning-rate ratio of 0.2, eight semantic units with slot competition enabled, and a fixed joint-prediction checkpoint supplying the encoder initialisation. The checkpoint reported in the main text is selected on dev; test is traversed once, at final evaluation, and is not used to select a checkpoint.

Per-domain metrics. Per-domain figures are recomputed from the stored predictions rather than read from the trainer, and reconcile with the trainer’s own pooled figures to within 5×10^{-5} on all six metrics.

D.2 PARAMETER SENSITIVITY

Each sweep varies one parameter from the default configuration of Table 2 and holds the rest fixed: the upstream pair and joint products, the encoder initialisation, the evaluation files and the seed. Every point is a single run at seed 173, except for the seed sweep below.




Figure 2: Sensitivity to the evidence budget K in the three domains. The default is $K=5$, the checkpoint of Table 2. K is a padding dimension, so the text reads the curves against the priors each setting consumes.

Evidence budget: the ideas run out of priors before the model runs out of capacity. Figure 2 varies K , the number of priors each idea contributes, over $K = 2, 3, 4, 5, 6, 8$. Accuracy rises steeply from $K=2$ to $K=4$ (0.731 $\rightarrow$ 0.806 on Graph ML, 0.716 $\rightarrow$ 0.772 on computer vision, 0.726 $\rightarrow$ 0.789 on time series) and then flattens, with the same knee in all three domains. The flattening comes from the evidence rather than from the model. Because every sample is padded to exactly K slots, the priors a setting consumes stop growing once K exceeds the candidates an idea carries, and the knee sits at three to four consumed priors. Raising K from 5 to 8 adds less than one consumed prior, because about a third of the ideas carry six or more candidates and fewer than one in seven carry eight or more.

Training-set size: more data helps across the board. Figure 3 varies the number of training ideas over 25%, 50%, 75% and 100% of the training split, drawn as nested stratified prefixes so that consecutive points add samples without reordering what came before. Macro-F1 rises steadily with the fraction on Graph ML (0.7627 $\rightarrow$ 0.8010) and time series (0.7573 $\rightarrow$ 0.7760). Computer vision




Figure 3: Sensitivity to the amount of EGRD training data in the three domains. Development and test sets stay fixed, and 100% is the default.

levels off after 50% and ends at 0.7669 against a peak of 0.7788, a gap of 0.012. More training data helps in most of the domains we measure, and the spread across the four fractions stays small: accuracy moves by at most 0.039 in any domain.




Figure 4: Sensitivity to encoder input length on the three domain test splits. The default is 384, and every setting scores the same test ideas and labels.

Encoder input length: 128 tokens already cover the comparison. Figure 4 varies the encoder input length over 128, 192, 256 and 384 tokens. The four lengths separate the domains by little: no domain moves monotonically with input length, and tripling the input moves Macro-F1 by at most 0.017 in any domain. At 128 tokens the encoder already covers the title-and-abstract comparison for most target–prior pairs, so a longer window reads much the same text and the curves stay flat. QWK is the flattest panel, spanning 0.0076 on Graph ML, 0.0023 on computer vision and 0.0062 on time series across the four lengths, so the ordering the model produces is close to invariant to how much of the text it reads. The flatness is consistent with the decomposition reading a target and its priors through short unit-level comparisons rather than through long spans of text.

Table 6: Per-domain seed variability. Entries are sample standard deviations across seeds 173, 174 and 175 with identical data, code and configuration.

Domain	Acc	Ma-F1	QWK	MAE	High-F1
Graph ML	0.0066	0.0050	0.0048	0.0067	0.0022
Computer vision	0.0040	0.0021	0.0032	0.0027	0.0037
Time series	0.0050	0.0045	0.0031	0.0051	0.0069

Seed stability: the reported model is not a selected seed. Table 6 varies nothing but the random seed, running the default configuration on seeds 173, 174 and 175 with identical data and code




Figure 5: Left: the eleven contribution types, placed by classical multidimensional scaling on the cosine distances between their centroids in the fused representation; two dimensions carry 61.9% of the centroid spread, and point area is class frequency. Colour is the seven-way coarse grouping the model is supervised with, so same-coloured points sitting apart means the layout is not a replay of that supervision; the dashed segments mark the two closest pairs in the whole map, both of which cross groups. Centre: each of the 55 type pairs, its centroid cosine against how often the model swaps the two. Right: held-out accuracy of a linear probe for the eleven-way type, averaged over three seeds.

and reporting the per-domain sample standard deviations. The spread is 0.004–0.007 in accuracy and 0.002–0.005 in Macro-F1, so the headline numbers describe the configuration rather than one fortunate run, and the three sweeps above should be read by the shape of a curve rather than by the difference between adjacent points.

E WHAT THE REPRESENTATION ORGANISES

The grade our model predicts and the representation it predicts it from are not the same object, and the second is the more informative one to inspect. This section reads six stage representations out of a single forward pass of one frozen checkpoint over the 3,352 test ideas: the target encoding, the covered pathway z_C , the two residual pathways z_U and z_I , the global feature context, and the fused vector the prediction heads receive.

Choosing a target the input scalars cannot reach. Two quantities enter the model as features and between them carry most of the four-way novelty grade: the joint coverage of the retrieved priors and the number of admissible candidates. A multinomial probe on those two scalars and their low-order interactions reaches 0.681 accuracy on the grade, against a 0.329 majority baseline and the model’s own 0.778. Any structure we find by probing for the grade would therefore be reporting, in large part, on those two numbers.

The residual type does not have this problem. It is an eleven-way label naming what kind of contribution remains once the priors are accounted for — a new mechanism, a non-trivial combination, a minor variant, a new evaluation protocol, and so on — and it is not a function of how much coverage there is. The same two-scalar probe reaches only 0.368 on it against a 0.275 majority, while the model’s own head reaches 0.651. The gap of twenty-eight points is the part that had to be read out of the text, and it is what the rest of this section examines. Figure 5 collects the result.

The types are laid out, and the layout predicts the model’s own errors. Projecting the eleven centroids to two dimensions recovers 61.9% of their spread, so the arrangement is close to planar. Its first axis correlates at -0.905 with the mean novelty grade of each type across the eleven classes, which places the conceptual and task-defining contributions at one end and minor variants and recombinations at the other.

The arrangement is not decoration: it predicts where the model goes wrong. Over the 55 type pairs, the cosine between two centroids correlates at $r = 0.469$ with how often the model swaps those

two types. The model never mistakes a new data distribution for a new theoretical concept, the most distant pair in the map, and its single most frequent confusion, minor variant against non-trivial combination, is one of its closest.

That correlation could be a restatement of the novelty axis, since types of similar novelty would be confusable for that reason alone. It is not. Partialling out the difference in mean novelty between the two types raises the correlation to 0.600 rather than removing it. The clearest case is NEW_MECHANISM against NEW_TRAINING_SIGNAL, whose mean grades differ by 0.05 and which novelty therefore cannot separate at all: they are the closest pair in the entire map at cosine 0.951, and among the novelty-matched pairs they are also the most often confused.

The layout disagrees with the grouping the model was taught. Training supervises a seven-way coarse grouping alongside the fine label, so the layout could be that grouping made visible. It is not. Pairs inside a coarse group sit at mean cosine 0.783 and pairs across groups at 0.709, a separation far smaller than the spread within either set, and the four closest pairs in the map all cross groups. The model places MINOR_VARIANT beside NONTRIVIAL_COMBINATION at 0.902 although the supervision assigns them to different coarse types, and it pulls NEW_MECHANISM away from NEW_SYSTEM_DESIGN to 0.798 although the supervision groups them together. Both departures read as the more natural organisation: a minor variant and a recombination are two degrees of reusing what exists, whereas a new mechanism and a new system design are different kinds of contribution. We report this as an observation about the representation, not as evidence that the taxonomy should be redrawn; with only four same-group pairs available the comparison of the two means is descriptive.

Where the knowledge sits. Probing each stage separately for the eleven-way type puts the covered pathway lowest at 0.456, the target encoding at 0.546, the interaction residual at 0.554, the global context at 0.554, the unit residual highest among the single pathways at 0.568, and the fused vector at 0.586. Two readings follow. Every text-derived representation is well clear of the 0.368 the input scalars reach, so the contribution type is genuinely carried by the encoded text rather than by the retrieval statistics. And the target encoding alone already reaches 0.546, which says that much of the type is legible from how the idea is phrased — a target that says it combines two things announces its own category — so the evidence pathways add less here than the gap to the model’s 0.651 head might suggest.

F RINoBENCH EVIDENCE BUDGET

The blockwise rule lets the frozen checkpoint, trained with at most five candidate priors per idea, read more evidence than it ever saw in training by running once per block of five and taking the minimum. Raising the budget from 25 priors to 60 moves accuracy from 0.4116 to 0.4152 and QWK from 0.2182 to 0.2245. The gain is real but small, and it saturates: RINoBench ideas carry a median of 23 candidate priors, so beyond roughly 45 the later blocks cover too few ideas to matter. A thirteenth block was not run because only seven test ideas still carry a prior at that depth and the join against the joint-coverage stage returns empty.