




Eleven contribution types, placed by how the model represents them; colour is the grouping it was taught



A regularity the model did learn: the shape of the contribution, not its grade
Closer types are confused more
 $r = +0.469, +0.600$ controlling for novelty



The input scalars cannot do this task — the text representations can



3,816 test ideas, one frozen checkpoint, one forward pass. Linear probes are held-out and averaged over three seeds; the model's own head is its reported prediction.